



# 第7讲 一致性和复制




§7.1 复制与副本管理

§7.2 副本一致性模型

§7.3 副本一致性协议


# §7.1 复制与副本管理

- Replication:
  - 数据或者资源部署多个副本
  - 每个副本都向客户端服务
    - 与备份不同
- 作用:
  - For Reliability
  - For Scalability
    - Number, location



# Major Issues in Replication

- Replica placement
  - 何处、何时、由谁来负责副本
  - Tradeoff between access benefit and update cost
- Consistency among replicas
  - 何种机制来保持副本的一致性
  - 确保冲突的操作在所有副本按照相同的顺序执行
    - 读写冲突 (Read-write conflict) , 读操作和写操作并发执行
    - 写写冲突 (Write-write conflict) , 两个并发的写操作



# Replica Management


- To decide where, when, and by whom replicas should be placed
  - Replica Server Placement
    - Finding the best locations for replica servers
  - Data Content Placement
    - Finding the best servers for placing content
- To distribute content to replicas

# Replica Server Placement

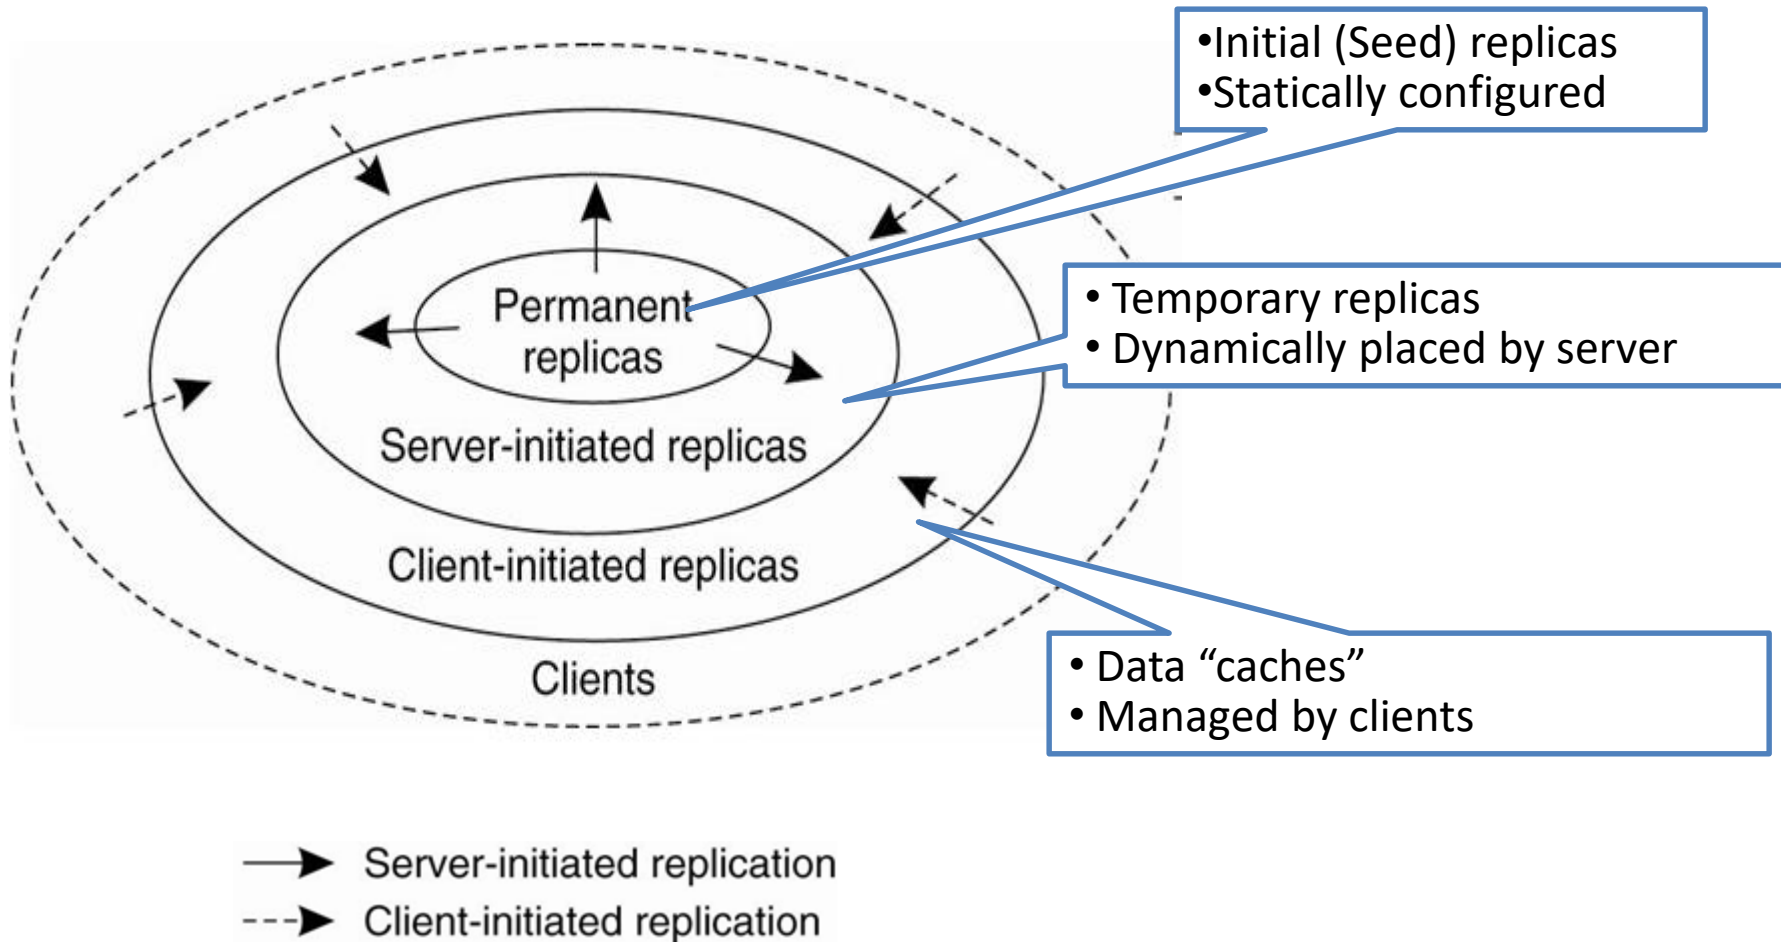
- $K$  out of  $N$  locations
- Client-aware method: high complexity
  - Objective function: average distance between clients and servers
- Client-unaware method: lower than previous, but still high complexity
  - Assuming clients are uniformly distributed
  - Greedy: choose routers with the largest links
- Region-based method: low complexity
  - A region is identified to be a collection of nodes accessing the same content, but for which the internode latency is low.

# Region-based method

- The entire space is partitioned into cells.
- The  $K$  most dense cells are then chosen.




# Data Content Placement



# Server-initiated Replicas


- Two considerations
  - To reduce the load of servers – near server
  - To reduce the access delay at clients – near client





# Server-initiated Replicas

- 对来自不同客户端的请求计数



- 记录每个文件/数据的访问次数，当作来自最靠近客户端服务器的请求；
- 如果请求数量低于阈值 $D$  -> 删除文件；
- 如果请求数量超过阈值 $R$  -> 复制文件；
- 如果请求数量在  $D$  和  $R$  之间 -> 移动文件；

# Client-initiated Replicas

- Client determine what to cache
- Servers may be involved for consistency
- Caches may be shared by more than one clients

# Content Distribution

- To propagate update to replica servers
- What to distribute
  1. Propagate only a **notification** of an update.
  2. Transfer **data** from one copy to another.
  3. Propagate the update **operation** to other copies
- 注意

没有哪一个方法是最佳的选择，高度依赖于可用的网络带宽和副本上的读写比率

# Content Distribution

- How to distribute
  - Push: suitable for permanent/server-initiated replicas
    - High read-to-update ratio
  - Pull: suitable for caches
    - Low read-to-update ratio

Issue	Push-based	Pull-based
State at server	List of client replicas and caches	None
Messages sent	Update (and possibly fetch update later)	Poll and update
Response time at client	Immediate (or fetch-update time)	Fetch-update time

# Content Distribution

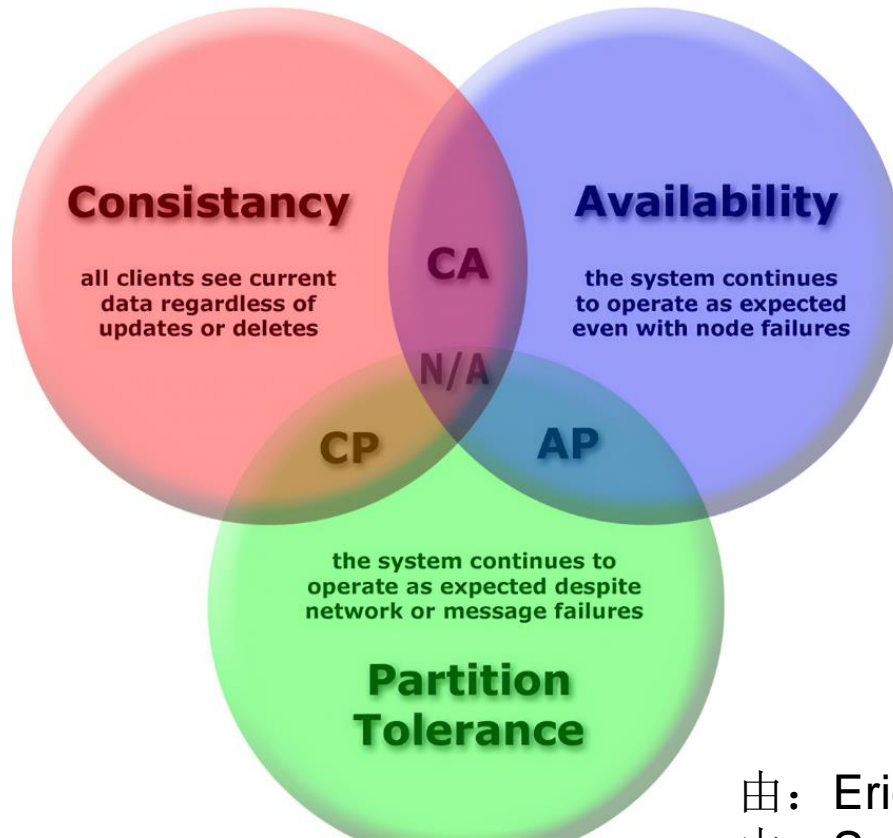
- How to distribute:
    - Lease-based: hybrid of push + pull
      - Lease (租约) 是服务器所作的承诺
      - 在lease指定的时间内服务器会把更新推给客户
      - lease到期则需要客户端通过pull方式更新
    - 关键问题：确定租约期限
- **Age-based leases**: An object that hasn't changed for a long time, will not change in the near future, so provide a long-lasting lease
  - **Renewal-frequency based leases**: The more often a client requests a specific object, the longer the expiration time for that client (for that object) will be
  - **State-based leases**: The more loaded a server is, the shorter the expiration times become

## §7.2 Consistency Models

- Consistency Model
  - A contract between processes and the data store.
  - It says that: if processes obey certain *rules*, the store promises to work *correctly*.
- “Correctness” :
  - If a process performs a read operation on a data item, the operation should return a value that shows the results of *the last write operation* on that data.
- Which is the last write, without global clock?
  - Consistency model defines the values that a read operation can return.
  - Tradeoff between *consistency level* and *maintenance cost*

# CAP理论

- 指一个分布式系统中 CAP三者不可得兼



- ①一致性：客户端的读操作要么读到最新的数据，要么读取失败。
- ②可用性：任何客户端的请求都能得到响应数据。
- ③分区容忍性：当消息丢失或延迟到达时，系统仍会继续提供服务，不会挂掉。



由：Eric A. Brewer 在PODC 2000特邀报告中提出  
由：Seth Gilbert, **Nancy Lynch** 正式证明

# Consistency Models

- Data-centric consistency models
  - Continuous consistency (on data content)
  - Update order consistency (on data operations)
    - Sequential consistency
    - Causal consistency
    - Grouping operations
- Client-centric consistency models
  - Monotonic reads
  - Monotonic writes
  - Read your writes
  - Writes follow reads



# Continuous Consistency

- Defining the degree of “inconsistency”
- Deviation in numerical values between replicas
  - 如：股票价格，可以规定差值不超过0.01元，或0.5%
- Deviation in staleness between replicas
  - 如：天气预报数据，可以规定差别不超过1小时
- Deviation with respect to the ordering of update operations
  - 本地更新但未达成全局一致的那些临时操作的数量
  - 如：可以规定允许3个或5个操作为临时性，有可能需要回滚重新执行

# Consistency Unit

- Conit: the unit of in consistency maintenance.
- 可以用于定义一致性水平; Conit大小影响数据更新成本
- 如: 3个值 (g,p,d) 设置为一个Conit
  - 顺序偏差: 未被另一副本提交的操作数量
  - 数值偏差 = (未接收到的更新次数, 偏差权重)
    - 偏差权重 = 已提交的值与未收到的操作产生的结果之间的最大差值

Replica A

Conit		
	d = 558	// distance
	g = 95	// gas
	p = 78	// price
Operation		
< 5, B>	g ← g + 45	[g = 45]
< 8, A>	g ← g + 50	[g = 95]
< 9, A>	p ← p + 78	[p = 78]
<10, A>	d ← d + 558	[d = 558]

Vector clock A = (11, 5)  
 Order deviation = 3  
 Numerical deviation = (2, 482)

Replica B

Conit		
	d = 412	// distance
	g = 45	// gas
	p = 70	// price
Operation		
< 5, B>	g ← g + 45	[g = 45]
< 6, B>	p ← p + 70	[p = 70]
< 7, B>	d ← d + 412	[d = 412]

Vector clock B = (0, 8)  
 Order deviation = 1  
 Numerical deviation = (3, 686)

确定提交过的操作。

**Figure 7.2:** An example of keeping track of consistency deviations.

# Sequential Consistency

- Notations

P1:	W(x)a	
<hr/>		
P2:	R(x)NIL	R(x)a

- Definition: *the result of any execution is the same as if*
  - *the (read and write) operations by all processes on the data store were executed in same sequential order and*
  - *the operations of each individual process appear in this sequence in the order specified by its program*

# Sequential Consistency

- No “time” in the definition of sequential consistency model
- A operation sequence is *valid* provided that *all processes see the same sequence*.

P1:	W(x)a		
P2:	W(x)b		
P3:		R(x)b	R(x)a
P4:		R(x)b	R(x)a

(a)

P1:	W(x)a		
P2:	W(x)b		
P3:		R(x)b	R(x)a
P4:		R(x)a	R(x)b

(b)

Figure 7-5. (a) A sequentially consistent data store.  
(b) A NOT sequentially consistent data store.

# Sequential Consistency

Process P1	Process P2	Process P3
$x \leftarrow 1;$ $\text{print}(y, z);$	$y \leftarrow 1;$ $\text{print}(x, z);$	$z \leftarrow 1;$ $\text{print}(x, y);$

Figure 7-6. Three concurrently-executing processes.

$x \leftarrow 1;$ $\text{print}(y, z);$ $y \leftarrow 1;$ $\text{print}(x, z);$ $z \leftarrow 1;$ $\text{print}(x, y);$  Prints: 001011 Signature: 001011	$x \leftarrow 1;$ $y \leftarrow 1;$ $\text{print}(x, z);$ $\text{print}(y, z);$ $z \leftarrow 1;$ $\text{print}(x, y);$  Prints: 101011 Signature: 101011	$y \leftarrow 1;$ $z \leftarrow 1;$ $\text{print}(x, y);$ $\text{print}(x, z);$ $x \leftarrow 1;$ $\text{print}(y, z);$  Prints: 010111 Signature: 110101	$y \leftarrow 1;$ $x \leftarrow 1;$ $z \leftarrow 1;$ $\text{print}(x, z);$ $\text{print}(y, z);$ $\text{print}(x, y);$  Prints: 111111 Signature: 111111
(a)	(b)	(c)	(d)

Figure 7-7. Four valid execution sequences for the processes of Fig. 7-6. The vertical axis is time.

Totally 720 (6!) sequences; 90 of them are valid.

# Causal Consistency

- Definition
  - *Writes that are potentially causally related must be seen by all processes in the same order.*
  - *Concurrent writes may be seen in a different order on different machines.*
- Recall “Causality”
  - If event *b* is *caused* or *influenced* by an earlier event *a*,
  - Everyone else should first see *a*, then see *b*.

Weaker than sequential consistency

# Causal Consistency

P1:	W(x)a			W(x)c
P2:		R(x)a	W(x)b	
P3:		R(x)a		R(x)c
P4:		R(x)a		R(x)b

Figure 7-8. A causally-consistent sequence but not sequentially consistent

# Causal Consistency

P1: W(x)a

P2: R(x)a W(x)b

P3: R(x)b R(x)a

P4: R(x)a R(x)b

(a)

P1: W(x)a

P2: W(x)b

P3: R(x)b R(x)a

P4: R(x)a R(x)b

(b)

Is the operation sequence causally consistent?



# Grouping Operations

- Operation consistency with mutual exclusion mechanism

**ENTER\_CS**, R(), W(), R(),..., **LEAVE\_CS**

- The consistency granularity is higher
  - Critical section: a group of reads and writes.
- Operation:
  - Synchronization variables (locks)
  - Acquire → Read/Write → Release

# Grouping Operations

- Acquiring a lock can succeed only when all updates to its associated shared data have completed.
  - 在一个进程对被保护的共享数据的所有更新操作执行完之前，不允许另一个进程执行对同步化变量的获取访问。
- Exclusive access to a lock can succeed only if no other process has exclusive or nonexclusive access to that lock.
  - 在更新一个共享数据项之前，进程必须以互斥模式进入临界区，以确保不会有其他进程试图同时更新该共享数据。
- Nonexclusive access to a lock is allowed only if any previous exclusive access has been completed, including updates on the lock' s associated data.
  - 如果一个进程要以非互斥模式进入临界区，必须确保临界区获得了被保护共享数据的最新副本。

P1: Acq(Lx) W(x)a Acq(Ly) W(y)b Rel(Lx) Rel(Ly)

---

P2: Acq(Lx) R(x)a R(y) NIL

---

P3: Acq(Ly) R(y)b

Figure 7-10. A valid event sequence for entry consistency

# Client-centric Consistency Models

- Weaker than data-centric ones
- Assuming restricted concurrency
  - E.g.
    - A database may be rarely updated
      - Few write-write conflicts
    - A database may be updated by only a special process
      - No write-write conflicts
    - Users may allow a quite high degree of inconsistency
      - E.g. web pages
- Guarantee *eventual consistency*

# Eventual Consistency

- Suitable for
  - (large-scale) distributed and replicated databases that tolerate a relatively high degree of inconsistency.
- Key point:
  - If no updates take place for a long time, all replicas will gradually become consistent.  
( In the absence of updates, all replicas converge toward identical copies of each other. )

Eventual consistency essentially requires only that:  
updates are guaranteed to propagate to all replicas.

# Eventual Consistency

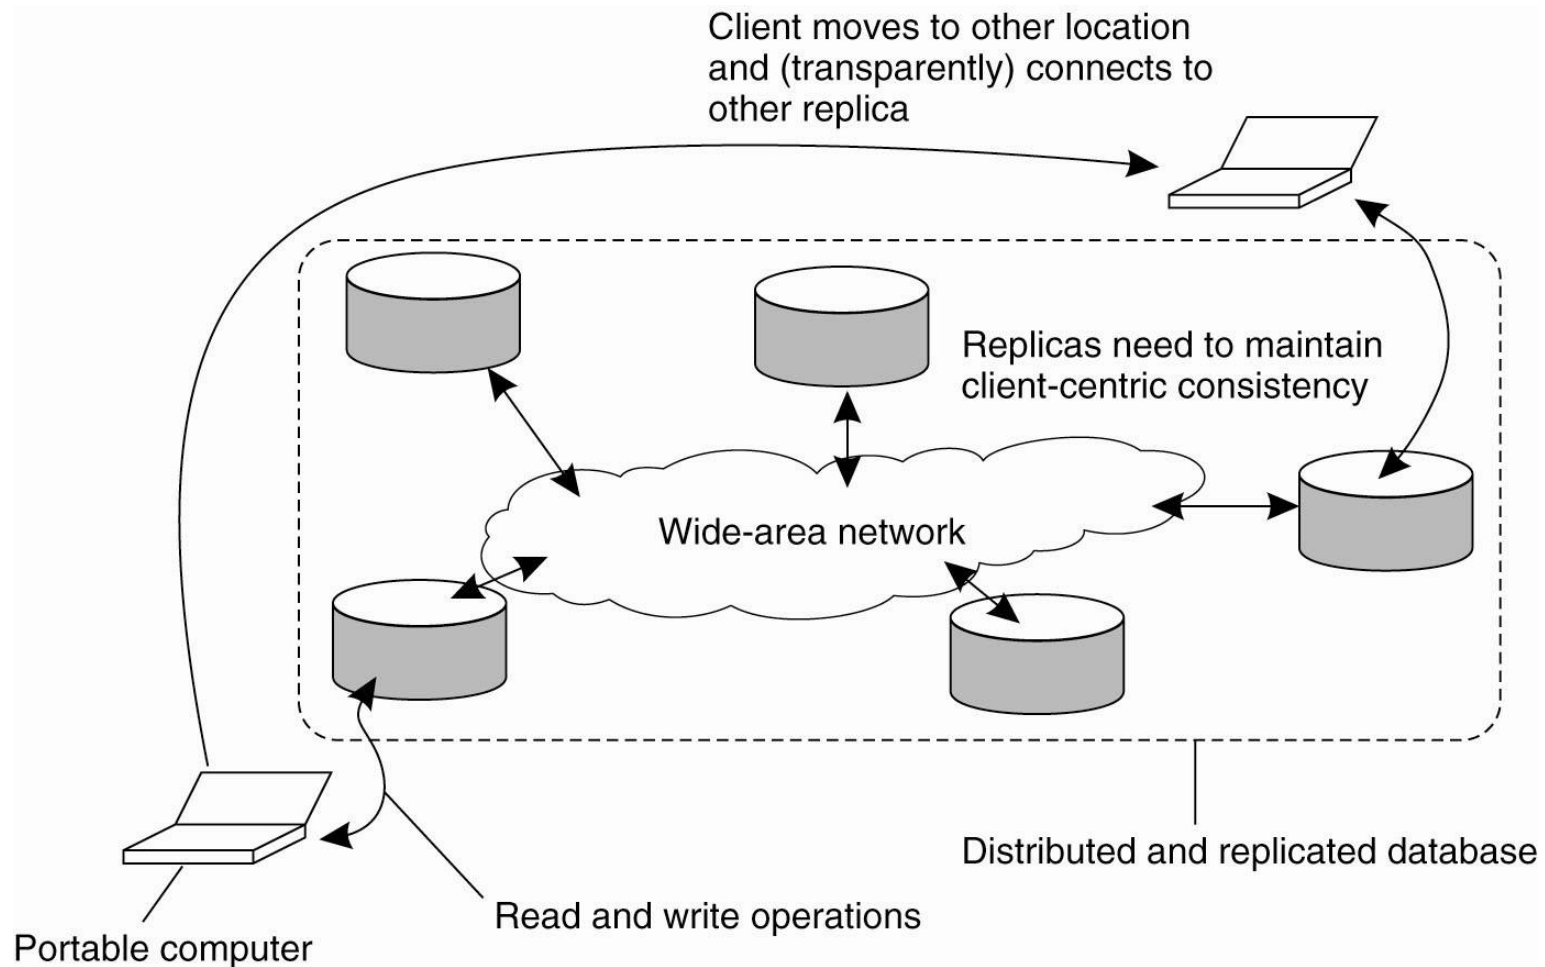


Figure 7-11. Inconsistency in eventually consistent database

# Client-centric Consistency

- To avoid inconsistency in eventually consistent systems
- Key points:
  - Provides **guarantees** for a **single client** concerning the consistency of **accesses** to a data store **by that client**.
  - **No guarantees** are given concerning concurrent accesses **by different clients**.
- Four Models
  - Monotonic reads
  - Monotonic writes
  - Read your writes
  - Writes follow reads



# Monotonic Reads 单调读



- If a process reads the value of a data item  $x$ , then any successive read operation on  $x$  **by that process** will always return that same value or a more recent value.

( if a process has seen a value of  $x$  at time  $t$ , it will **never see an older** version of  $x$  later.)

E.g. mail service

“读到值只会是越来越新”

# Monotonic Reads



$L_i$ : 本地副本i

$W_i(x_n)$ : 进程i对x写为版本n

$R_i(x_n)$ : 进程i读到x版本n

$W_i(x_m;x_n)$ : 在完成版本m之后，进程i写入x的版本n

$W_i(x_m|x_n)$ : 进程i并发写入x的两个版本版本m和n

Fig. 7-16: (a): Guaranteed (b): Not.





# Monotonic Writes 单调写

- A write operation by a process on a data item  $x$  is completed before any **successive write** operation on  $x$  by the **same process**.

(一个进程对数据项  $x$  执行的写操作必须在该进程对  $x$  执行任何后续写操作之前完成。

If need be, the new write must wait for old ones to finish.)

E.g. software library

“自己提供的新版本不会被旧版本覆盖”

# Monotonic Writes

$$\begin{array}{l} \text{L1: } W_1(x_1) \\ \hline \text{L2: } W_2(x_1; x_2) \quad W_1(x_2; x_3) \end{array}$$

(a)

$$\begin{array}{l} \text{L1: } W_1(x_1) \\ \hline \text{L2: } W_2(x_1 | x_2) \quad W_1(x_1 | x_3) \end{array}$$

(b)

$$\begin{array}{l} \text{L1: } W_1(x_1) \\ \hline \text{L2: } W_2(x_1 | x_2) \quad W_1(x_2; x_3) \end{array}$$

(c)

$$\begin{array}{l} \text{L1: } W_1(x_1) \\ \hline \text{L2: } W_2(x_1 | x_2) \quad W_1(x_1; x_3) \end{array}$$

(d)

Fig. 7-17: (a) yes. (b) no. (c) no. (d) yes (although  $x_1$  has apparently overwritten  $x_2$ ).



# Read Your Writes 读写一致性

- The **effect** of a **write** operation by a process on data item  $x$  will always be **seen by** a **successive read** operation on  $x$  by the same process.

(一个写操作总是在同一进程执行的后续读操作之前完成，而不管这个后续读操作发生在什么位置。)

E.g. Password changing

“自己总是看到自己写过的最新版本”

# Read Your Writes

$$\begin{array}{lcl} \text{L1:} & W_1(x_1) & \\ \hline \text{L2:} & W_2(x_1; x_2) & R_1(x_2) \end{array}$$

(a)

$$\begin{array}{lcl} \text{L1:} & W_1(x_1) & \\ \hline \text{L2:} & W_2(x_1 | x_2) & R_1(x_2) \end{array}$$

(b)

Fig. 7-18: (a): Guaranteed (b): Not.

# 读写一致性例子

- 更新Web页面，并且保证Web浏览器能够展示最新的版本的数据，而不是缓存的内容；

**facebook**

邮箱或手机号

密码

登录

忘记帐户?

联系你我，分享生活，尽在 Facebook



**注册**

永久免费使用

姓

名

手机号或邮箱

创建密码



# Writes Follow Reads 写读一致性

- A **write** operation by a process on a data item  $x$  following a **previous read** operation on  $x$  by the same process is guaranteed to take place on the **same or a more recent value** of  $x$  that was read.

(进程对数据项  $x$  所执行的任何后续写操作都会在  $x$  的副本上执行，此时该副本已经具有该进程最近读取的或更新版本的值。)

E.g. bbs, newsgroup

“所见到过的必须已经出现”

# Writes Follow Reads

$$\begin{array}{c}
 \text{L1:} \quad W_1(x_1) \quad R_2(x_1) \\
 \hline
 \text{L2:} \quad W_3(x_1; x_2) \quad W_2(x_2; x_3)
 \end{array}$$

(a)

$$\begin{array}{c}
 \text{L1:} \quad W_1(x_1) \quad R_2(x_1) \\
 \hline
 \text{L2:} \quad W_3(x_1 | x_2) \quad W_2(x_1 | x_3)
 \end{array}$$

(b)

Fig. 7-19: (a): Guaranteed (b): Not.

## §7.3 Consistency Protocols

- Data-centric consistency models
  - Continuous consistency (Data content consistency)
  - Update order consistency
    - Sequential consistency
    - Causal consistency
    - Grouping operations
- Client-centric consistency models
  - Monotonic reads
  - Monotonic writes
  - Read your writes
  - Writes follow reads



# 连续一致性：限定数值偏差

- Every server  $S_i$  has a log, denoted as  $L_i$ .
- Consider a data item  $x$  and let  $val(W)$  denote the numerical change in its value after a write operation  $W$ . Assume that

$$\forall W : val(W) > 0$$

- $W$  is initially forwarded to one of the  $N$  replicas, denoted as  $origin(W)$ .  $TW[i, j]$  are the writes executed by server  $S_j$  that originated from  $S_i$ :

$$TW[i, j] = \sum \{ val(W) | origin(W) = S_j \ \& \ W \in L_i \}$$

Note

Actual value  $v(t)$  of  $x$ :

$$v(t) = v_{init} + \sum_{k=1}^N TW[k, k]$$

value  $v_i$  of  $x$  at server  $S_i$ :

$$v_i = v_{init} + \sum_{k=1}^N TW[i, k]$$

# 连续一致性：限定数值偏差

## Problem

We need to ensure that  $v(t) - v_i < \delta_i$  for every server  $S_i$ .

## Approach

Let every server  $S_k$  maintain a **view**  $TW_k[i, j]$  of what it believes is the value of  $TW[i, j]$ . This information can be **gossiped** when an update is propagated.

## Note

$$0 \leq TW_k[i, j] \leq TW[i, j] \leq TW[j, j]$$

- 基本操作
  - $S_i$  扩散写操作 originating from  $S_j$  to  $S_k$  (写操作传播机制)
  - $S_k$  会获得  $TW[i, j]$ , 如发现更新操作步调不一致, 把写操作从日志中转发给  $S_i$
  - 转发操作可以把  $S_k$  的视图  $TW_k[i, k]$  往  $TW[i, k]$  靠近
  - 当应用程序提交一个新的写操作时,  $S_k$  会把其视图往  $TW[k, k]$  推, 从而导致
$$TW[k, k] - TW_k[i, k] > \frac{\delta_i}{N-1}$$
  - 但是本方法能够确保  $TW[i, k] - TW_k[i, k] \leq \delta_i$



# 连续一致性：限定陈旧度、顺序偏差

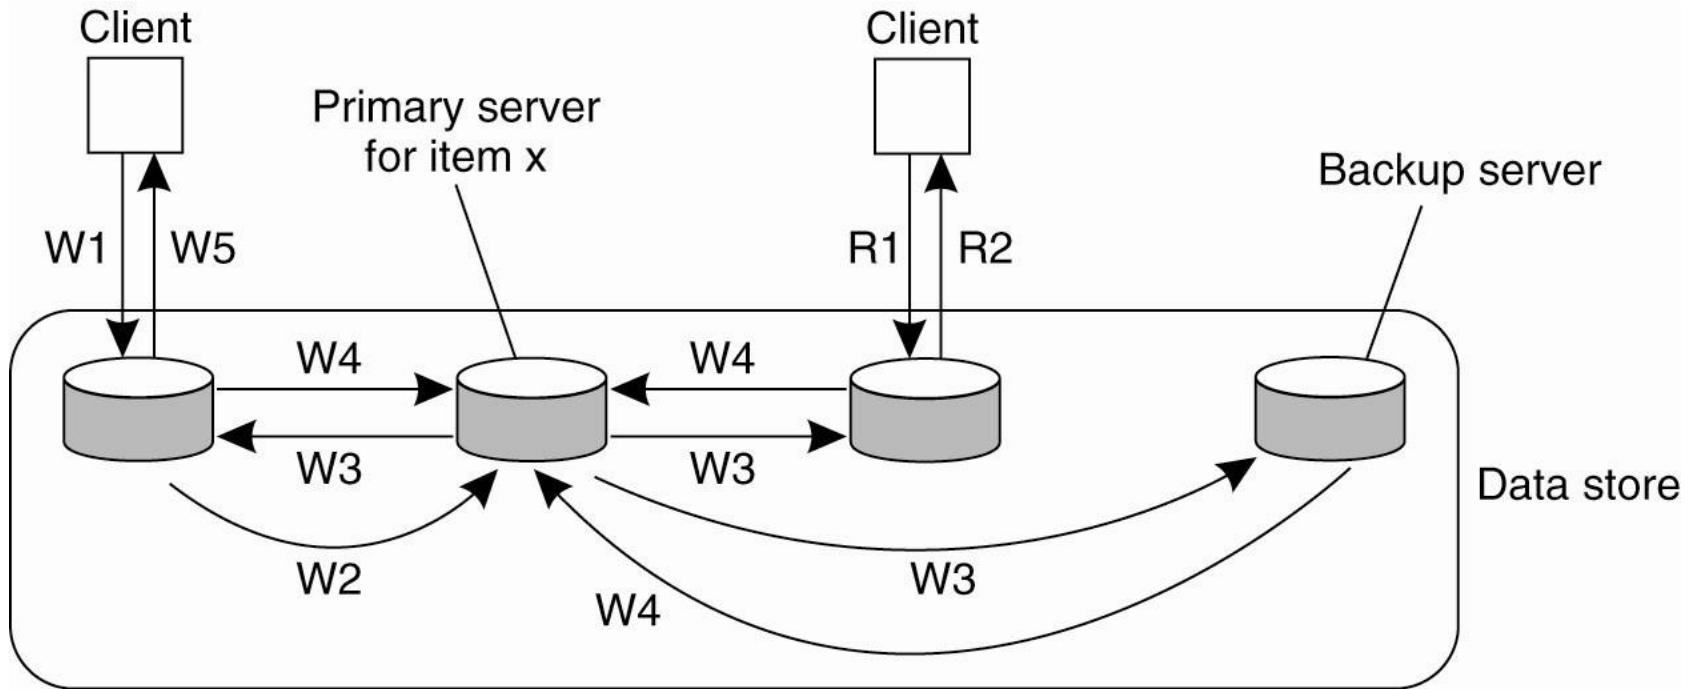
- 限定陈旧度
  - 服务器  $S_k$  保持实时向量时钟  $RVC_k$ 
    - $RVC_k[i] = T(i)$  为到时间  $T(i)$  时,  $S_k$  看到了已提交给  $S_i$  的所有写操作
  - 只要服务器  $S_k$  发现  $T[k] - RVC_k[i]$  将超出指定界限, 那么就拉入来自  $S_i$  的时间戳晚于  $RVC_k[i]$  的写操作
- 限定顺序偏差
  - 暂存写操作到本地队列
  - 当本地写队列的长度超过限时, 不再接受任何新提交的写操作, 按照相应的顺序提交写操作



# 操作一致性协议

- Primary-based Protocols
  - Remote-write
  - Local-write
- Replicated-write Protocols
  - Active replication
  - Quorum-based

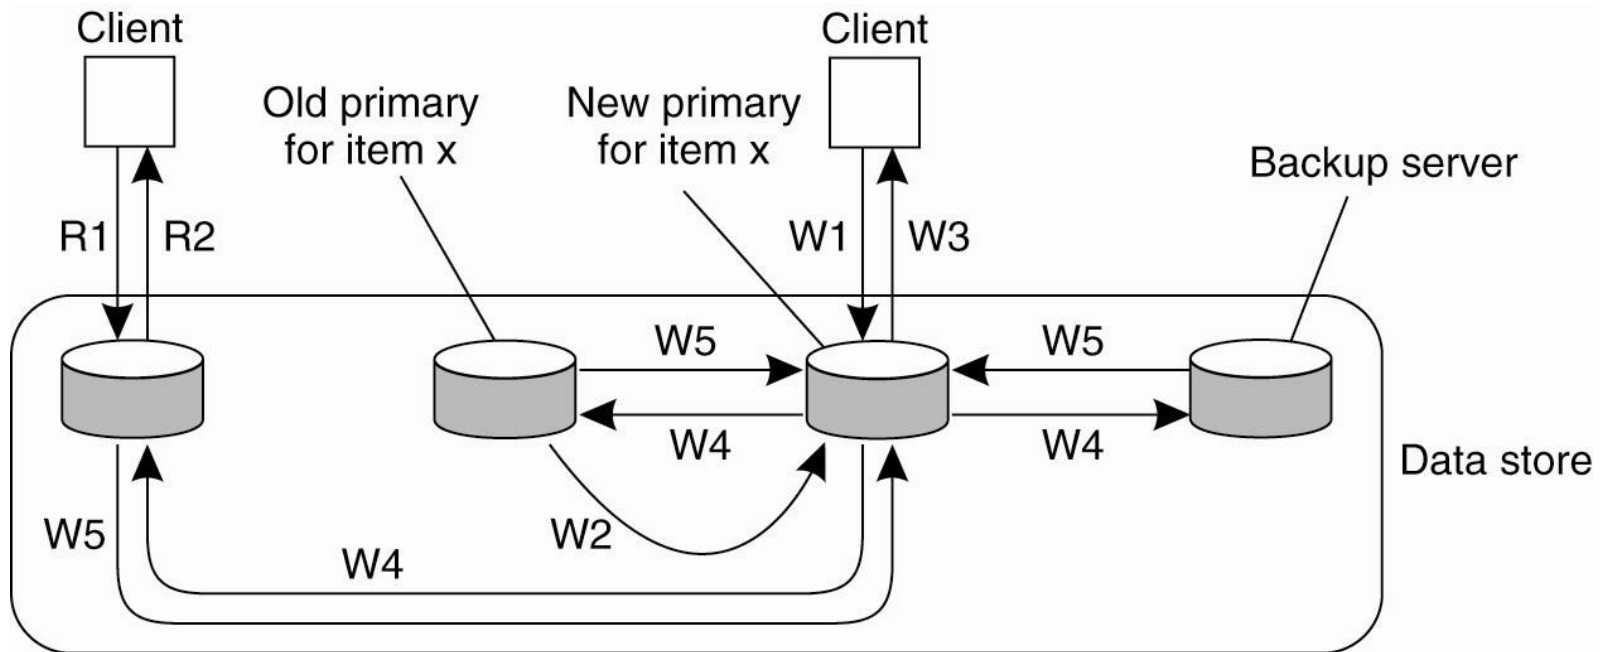
# 主副本协议: Remote-write Protocol



W1. Write request  
W2. Forward request to primary  
W3. Tell backups to update  
W4. Acknowledge update  
W5. Acknowledge write completed

R1. Read request  
R2. Response to read

# 主副本协议: Local-write Protocol




W1. Write request  
W2. Move item x to new primary  
W3. Acknowledge write completed  
W4. Tell backups to update  
W5. Acknowledge update

R1. Read request  
R2. Response to read

# 复制写协议：Active Replication Protocol

- The operation is forwarded to all replicas
- Operations to be carried out in the same order everywhere
- Requires totally ordered multicasts using either Lamport timestamps (e.g.) or a central coordinator.





# 复制写协议：Quorum-Based Protocol

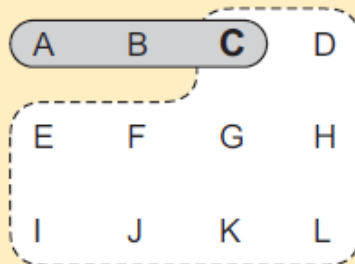
## Quorum-based protocols

Ensure that each operation is carried out in such a way that a majority vote is established: distinguish **read quorum** and **write quorum**


Three examples of the voting algorithm. (a) A correct choice of read and write set. (b) A choice that may lead to write-write conflicts. (c) A correct choice, known as ROWA (read one, write all)

1.  $N_R + N_W > N$

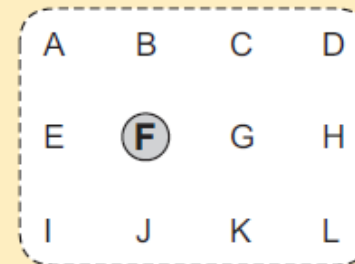
2.  $N_W > N/2$



$N_R = 3, N_W = 10$



$N_R = 7, N_W = 6$



$N_R = 1, N_W = 12$



# 客户为中心的一致性协议

- 客户检查自己的更新是否完成
  - 每个写操作被唯一标识;
  - 每个客户维护两个集合:
    - 读操作集{客户的读操作相关的写操作};
    - 写操作集{客户的写操作}
  - 单调读/写:
    - 相关的读/写操作集与读/写请求一起发送服务器;
    - 执行读操作前检查是否所有写已经在本地执行;
    - 没有的话: 联系其他服务器进行更新, 或转发读请求出去。
  - 读写一致 (写后读): 读之前检查写操作集。
  - 写读一致 (读后写): 写之前检查读操作集, 并将读操作集加入写操作集。

# A Summary

- Replica management
  - 副本/缓存类型; 副本放置、内容分发
- Consistency models
  - 数据为中心的一致性 (读写并重系统)
    - 连续 (持续) 一致性、操作一致性
  - 用户为中心的一致性 (读为主)
    - 单调读、单调写、读写一致 (写后读)、写读一致 (读后写)
- Consistency protocols
  - 数据为中心:
    - 持续一致性协议: 更新扩散过程限制偏差
    - 操作一致性协议: 主备份写、全复制写、多数 (Quorum) 写
  - 用户为中心: 读写集检查

# Homework Questions

1. 请分析讨论，与sequential consistency相比，eventual consistency的优势和价值，并通过例子进行说明。
2. 下面Causal consistency的操作例子，最后的两个读操作应该返回什么结果？

P1:	W(x)a		
P2:		R(x)a	W(y)b
P3:			R(y)b R(x)?
P4:			R(x)a R(y)?

3. 给出一个实现数据副本的因果一致性的方法思路。